

Investor Prospectus

Brand visibility for the AI search era

The Ask: \$1.75M at \$5M pre-money valuation

Structure: Priced equity round, approximately 26% post-money dilution

Runway: 18 months from initial closing (Seed target Q2 2028)

Closing: Rolling close — initial closing 31 December 2026 (minimum \$500K), final closing 31 March 2027

Use of funds: Three hires (tech, sales, customer success), channel investment, SOC 2 Type II audit, Delaware conversion

CONFIDENTIAL — August 2026

CiteMetrix LLC · Norwalk, Connecticut

Note on corporate form. The operating entity is CiteMetrix LLC, a Connecticut limited liability company. Conversion to a Delaware C-corporation will be completed at or before the closing of this round, as required to issue the Seed Preferred described in Section 11. Conversion costs are included in the Section 9 allocation.

investors@citemetrix.com · citemetrix.com/investors

This document is intended for accredited investors only and is not an offer to sell securities. See “Important notices” on the final page.

1. Executive summary

CiteMetrix is the AI-search-visibility platform built for the brand buyer. We monitor how nine AI platforms — ChatGPT, Claude, Perplexity, Gemini, Copilot, Mistral, Grok, DeepSeek, and Google AI Overview — talk about consumer brands, quantify that visibility with a defensible 0–100 ModelScore™, detect hallucinations and provide one-click remediation, and surface competitive positioning that exports cleanly to a board deck. The platform is live at **v4.46.0**, built solo across more than 570 releases with no outside capital.

The category is funded and crowded, but every funded competitor is selling to the SEO team. Profound has raised \$155M including a \$96M Series C at a \$1B valuation. AthenaHQ, Otterly, Peec, and Knowatoya fill out the segment. All of them position at SEO practitioners, not at the brand managers and CMOs whose boards are now asking “*what is ChatGPT saying about us?*” The buyer is wide open. We estimate a 3–9 month window before the incumbents reposition.

Our go-to-market is channel-led. Healthcare, hospitality, and consumer-brand agencies are already fielding the AI-visibility question from their clients; they can wrap CiteMetrix into existing website scopes as a billable service line, where one signed agency multiplies across their full client roster. That thesis is now being tested in the field: an agency relationship that began in May — a

healthcare digital agency already working with us through our consulting practice — has expanded from one client brand to seven, going to nine, with no salesperson involved.

We are raising **\$1.75M at a \$5M pre-money valuation** to fund three hires (senior backend/AI engineer, agency channel sales, customer success and onboarding), SOC 2 Type II certification, and the channel investment required to convert agency-channel interest into recurring revenue. Eighteen months of runway to Seed milestones.

Snapshot

Product status	Live at v4.46.0 · paid conversion opened May 1, 2026
Releases shipped	570 , solo, no outside capital
Built-in tools	10+
AI platforms	9 , extensible via our Platform Registry
Live citation observations	31,614 (as of Aug 24, 2026) (18,095 brand-cited)
Monitored queries / ModelScores computed	264 active / 1,302 cumulative
Agency pipeline	Anchor agency at 7 client brands in pilot, going to 9 — unconverted
Revenue	Pre-revenue. No paying customers to date. No sales hire made.
Founder & CEO	Eric Richmond (30+ years SEO and search)
Strategic advisor	Brian Cauley
Raise / pre-money / dilution	\$1.75M / \$5M / ~26%
Structure	Priced equity round

2. The market opportunity

Why AI search matters now

AI search is replacing meaningful portions of the traditional Google search experience. Google's own AI Overviews surface above organic results on a growing share of branded queries, and ChatGPT, Claude, and Perplexity have moved from novelty to embedded research workflow in under 24 months. Published survey work suggests roughly half of B2B buyers now begin product research inside an AI assistant rather than a search engine; we cite that directionally rather than as a precise figure, and the thesis does not depend on the exact number.

For consumer brands the implication is structural: the SERP is no longer the only — or even the primary — discovery surface. When a brand manager's board asks whether the company is being recommended, ignored, or hallucinated about by the AI tools their customers actually use, no traditional SEO tool can answer.

Category sizing and funding

The AEO ("Answer Engine Optimization") category has attracted more than \$300M in combined disclosed funding over the last 18 months. Profound alone raised \$155M, including a \$96M Series

C at \$1B led by Lightspeed, Sequoia, and Kleiner Perkins. More than 150 AEO-adjacent tools have launched in ten months.

This is the moment in a category's life when capital is plentiful but positioning is unfinished. Every funded incumbent has chosen to pitch SEO teams — because that is the buyer they know, the buyer their network knows, the buyer whose budget they understand. The brand-side buyer sits one floor up in the same building, controls a meaningfully larger budget, and has the AI-visibility question being asked of them by their board. They are unserved.

The brand buyer profile

Target buyer is the brand manager or CMO at a consumer-facing company with \$100M+ in revenue, where AI-driven brand mentions translate into measurable business impact. Healthcare, hospitality, consumer goods, financial services, and direct-to-consumer brands are the priority verticals. The buying motion is not direct — these buyers do not typically procure SaaS tools independently. They buy through agencies.

The bet is that agencies are the right wedge into the brand buyer, not direct sales. Section 6 sets out the evidence accumulated against that bet since May.

3. The product

Architecture overview

CiteMetrix is a multi-tenant SaaS platform built on PHP/WordPress (production at citemetrix.com), with a Lightsail-hosted PWA at app.citemetrix.com for mobile and voice access, and a Flask-based admin portal for operations. Data layer is MariaDB on Amazon RDS.

The platform implements a BYOK (bring-your-own-key) model: customers supply their own API keys for the AI platforms they monitor. Marginal infrastructure cost per customer approaches zero, and unit economics are structurally healthier than competitors who resell inference at a markup and absorb the variable cost.

Nine AI platforms are integrated for citation tracking: ChatGPT (OpenAI), Claude (Anthropic), Perplexity, Gemini (Google), Copilot (Microsoft), Mistral, Grok (xAI), DeepSeek, and Google AI Overview (via DataForSEO). New platforms are added through our Platform Registry as the landscape evolves; the system is not locked to today's nine.

Differentiated capabilities

The platform ships **10+ built-in tools**. Five capabilities carry the differentiation.

ModelScore™ — a defensible 0–100 quantification of AI visibility. Composite of **Mention Score (45%)**, **Brand Demand (20%)** via Google Search Console, **Authority Transfer (20%)** via GA4 or Adobe Analytics, and **Technical Readiness (15%)**. The score is the board-level metric: it is the number that goes into the slide that goes into the room. Methodology is published on a public page; weights and inputs are transparent and defensible to a CMO or a board director. An unconnected data source is excluded and disclosed rather than scored as a silent zero.

Hallucination detection with one-click remediation. Customers load verified brand facts — regulatory language, product specifications, claim language. The platform surfaces every AI

response that contradicts those facts, checked semantically so paraphrased errors are caught too. Surfaced issues feed the Content Engine, which generates updates to the customer's own web properties (llms.txt, schema, content) so the next AI crawl pulls corrected ground truth. The closed loop — detect, remediate, verify — is what no SEO-tool competitor has.

Answer Position Estimator (*new in 4.28.0*). Checks whether a page answers its target query within the **top 30% of its content** — the region from which 44.2% of LLM citations originate, per Zyppy's ranking-factor research. Returns a position score, the passage that answers the query, and a specific recommendation naming what occupies the top of the page instead. It is one of the few AI-citation factors a content team can act on the same afternoon, and no competitor audits it.

Competitors with no website (*new in 4.29.0*). Competitive tracking in this category is anchored on a rival's domain, which makes it structurally blind wherever the competition has no website to point at. In pharma that is routine: a prescription brand's largest competitor is frequently its own molecule, its therapeutic class, or a generic — none of which any company has an incentive to build a page for. CiteMetrix matches competitors by name and alias rather than by domain, so a molecule, a class, or a generic is tracked like any other rival. Category-level mentions are reported as co-citation and are counted only where the brand is absent, so a brand is never scored against a term describing itself.

Citation-Attributed Content Change Tracker (*new in 4.28.0*). Snapshots the main content of the most-cited pages weekly, detects material changes with navigation and boilerplate stripped out, summarizes what changed in plain language, and — where the data supports it — shows how that page's share of AI citations moved afterward, judged against the site-wide trend so a scan-volume swing cannot masquerade as a gain.

The pitch line is: *"You rewrote your pricing page on the 24th. ChatGPT started citing it three days later. Perplexity didn't. Here is exactly what changed."* It catches changes nobody logged, including work another team shipped without telling marketing. **It never claims causation, and it suppresses any performance verdict when the citation data is too thin to support one.** That restraint is the moat rather than a caveat: a number a CMO can defend in a board meeting is worth more than a number a dashboard invented.

Both new tools are built on a shared, boilerplate-immune content extractor — platform infrastructure, not one-off features.

Reporting and data access

Reporting is the layer that turns measurement into a deliverable, and it splits two ways.

A self-serve report builder produces ready-made report templates and fully custom reports as PDF, native editable PowerPoint, or on-screen HTML with an embedded analyst answering questions against live data. Agencies can upload their own branded PowerPoint with {{tokens}} where the data belongs, and CiteMetrix fills them with live client numbers under the agency's logo. Reports schedule and deliver automatically.

A read-only Data Export REST API feeds citations, ModelScore trends, competitors, and hallucination data into Looker Studio, Tableau, or any BI tool.

Together: a builder for the many, an API for the power users. The white-label path is what converts the platform from a tool subscription into a billable agency deliverable.

The transparency moat

We maintain a public methodology page documenting how ModelScore is computed, what each component measures, how rate normalization handles missing data, and — explicitly — what the platform cannot determine. It states plainly that SOC 2 controls are in place and that the third-party audit has not yet been completed. No incumbent publishes this level of detail.

We treat methodology transparency as a category-defining choice. It converts skeptical buyers, gives the press something to write about, and pre-empts the “black box” objection every other AI-visibility tool faces.

4. Go-to-market

Channel-led, agency beachhead

Our motion is channel-led with agencies as the beachhead, not direct sales. The reasoning is structural rather than tactical: brand-side buyers at \$100M+ consumer companies do not procure SaaS tools through outbound. They procure capabilities through the agencies who already run their website, their content programs, their paid media, their analytics. When a brand manager’s board asks the AI-visibility question, the answer that arrives is not a SaaS demo. It is the agency saying “we already have a service line for that.”

The three-stage motion: (1) CiteMetrix sells to the agency; (2) the agency embeds CiteMetrix into existing website and content engagements as a margin-positive line item; (3) the brand manager has the answer when the board asks. The agency keeps the relationship and the renewal. CiteMetrix scales through the agency’s client roster with effectively zero customer acquisition cost per end-brand.

Agency pricing starts at \$299/mo and the most common expected tier is \$599–999/mo — deliberately set below the threshold requiring senior sign-off at most agencies, which means a sales cycle measured in weeks rather than quarters.

How agency revenue actually scales

This is the mechanic that matters, and it is easy to model incorrectly.

Every Agency tier includes unlimited client domains. An agency does not pay more for adding a client. Tiers are separated by **query pool** — the number of monitored prompts the agency can allocate across its roster. Revenue therefore scales with **monitoring depth**, not with client count.

Worked example from the live pipeline: an agency running nine client brands at roughly twelve monitored queries each needs about 108 queries. That exceeds Agency Medium’s 75-query pool and lands on **Agency Large at \$999/mo — approximately \$12K ARR from a single agency relationship**, with headroom to roughly sixteen client brands before another upgrade.

The strategic consequence: growth inside an account comes from deeper monitoring per brand, not from adding logos. Land-and-expand runs on query depth.

The anchor relationship

In May 2026 the Managing Director of Digital Strategy and Delivery at a healthcare-focused digital agency — a relationship rooted in our consulting practice — told us:

“Whenever I’m talking with clients about GEO and content optimization they are all over it. We’ll be building it into all our website scopes.”

That agency began with one client brand. It has since expanded, without a salesperson involved, to **seven client brands in active pilot, going to nine**. Pilots run under our account during evaluation and transfer to the agency’s own account at conversion via the built-in account-transfer tool.

The agency has not converted to paid. What the relationship demonstrates is the multiplication mechanic the channel thesis predicts — an agency adding its own clients to the platform before anyone has asked for a purchase order. Converting it is a sales function that does not currently exist. That is what this round funds.

Beyond the beachhead

The agency channel is the first 12–18 months. Parallel motions in development: PR firms (similar dynamic, crisis and reputation use cases), enterprise communications departments (longer cycle, higher LTV), and direct-to-brand for the largest consumer accounts. The channel-led wedge funds the build-out of the others; this round does not depend on them.

5. Competitive landscape

Five funded or notable competitors define the category. Every one positions at SEO teams.

Profound — category leader by funding (\$155M total, \$96M Series C at \$1B led by Lightspeed, Sequoia, Kleiner Perkins). Built proprietary crawling infrastructure for Fortune 500 compliance. 700+ enterprise customers, 18 months old. SOC 2 Type II certified. Their pitch is to Fortune 500 SEO teams and their pricing reflects it. Profound’s positioning is not a mistake — it is a deliberate choice that opens the brand-side wedge for a focused competitor.

AthenaHQ — newer entrant, enterprise positioning. Tracks 8 AI platforms. Limited public funding information. Early-form hallucination detection, no remediation closed loop.

Otterly, Peec, Knowatoa — the SMB SEO segment. All sub-\$5M funded, targeting small-business SEO practitioners, tracking 3–6 AI platforms each, offering monitoring without remediation at \$29–\$49 entry points. None has a defensible composite score, a hallucination remediation pipeline, or a mobile and voice surface.

The white space

Plotting the category on two axes — SEO buyer ↔ Brand buyer × Surface monitoring ↔ Analytical depth — produces four quadrants. Profound and AthenaHQ occupy SEO/analytical. Otterly, Peec, and Knowatoa occupy SEO/surface. Brand/surface is empty. The fourth quadrant, **Brand buyer + analytical depth, is where CiteMetrix sits alone.**

The competitive set's SEO-first positioning is publicly committed — bios, conference talks, sales materials, customer case studies. Walking it back to address the brand buyer is expensive in both time and credibility. We estimate a 3–9 month window before the incumbents reposition.

6. Traction and milestones

Beat 1 — Beta validated the methodology on real data

Closed beta ran from late 2025 through April 30, 2026 and opened early access to **39 sign-ups**. Most were early-adopters sampling the platform; the meaningful validation came from two real deployments run through the platform during this period. The first, a specialty-materials manufacturer and long-standing consulting client, was tracked across 30–40 monitored queries over three months and generated the bulk of early citation volume — stress-testing the measurement engine against a live brand at scale (its share-of-voice outcome is detailed in the competitive results). The second, a prescription respiratory brand (Xopenex HFA), remains monitored today and is the platform's clearest regulatory proof point: as of Aug 24, 2026 it carries **636 critical, still-open accuracy records tied to a single verified fact** — four AI platforms (DeepSeek, Copilot, Grok, Mistral) independently stating an approved pediatric age older than the FDA label, which indicates the drug for children 4 years and older. Across all monitored facts the domain carries 1,991 critical open issues. Every error is captured verbatim and checked against the brand's own verified regulatory language.

This is verifiable proof that the measurement works against real brands with real regulatory exposure — grounded in the drug's FDA label and the platform's own verified facts.

Beat 2 — We deepened the product before scaling the channel

Paid conversion opened May 1, 2026. Rather than scale the channel immediately, we closed real product gaps first. Between May and August the platform moved from v4.4 to **v4.28** — solo, with no outside capital, across more than 500 total releases.

Shipped in that window: the full self-serve report builder with PDF, native PowerPoint and on-screen output; per-client white-label branding; domain-level credential isolation and per-member access control; a read-only data export API; and the two new moat tools described in Section 3.

The release cadence is the evidence, and it is independently checkable through the public changelog.

Beat 3 — The channel thesis is validating ahead of the sales motion

Anchor agency: healthcare-focused digital agency (an existing consulting-practice relationship), engaged May 2026, now at **seven client brands in active pilot, expanding to nine**. No salesperson, no marketing spend, nothing sold to them. Unconverted.

Second agency: early conversation, client proof-of-concept loaded, founder demonstration scheduled September.

Platform data depth today: 31,614 citation observations (as of Aug 24, 2026) (18,095 brand-cited) across all nine platforms, **264 active monitored queries, 1,302 ModelScores computed**,

224 verified brand facts, 69 competitors tracked. Real production data, not synthetic — the substrate every measurement claim in this document rests on.

Revenue status: pre-revenue. No paying customers to date. No sales hire has been made.

An agency multiplying across its own client roster before anyone asked for a purchase order is the clearest available evidence that the wedge works. The reason it has not converted is that the function which converts it does not exist yet. That is the gap this round closes, and it is a known cost rather than an unknown risk.

7. Financial projections

Pre-seed financial projections are notoriously made up. What follows is channel-economics arithmetic with explicit assumptions, presented as scenarios rather than forecasts. Investors should discount accordingly; we make no claim to precision.

Assumptions

Pricing (two parallel tracks, live and current). Business: Solo \$79/mo (1 domain, 10 queries, 1 seat), Starter \$199/mo (3 domains, 25 queries per domain, 3 seats), Professional \$499/mo (10 domains, 50 queries per domain, 10 seats), Enterprise from \$1,499/mo (25 domains, 200 queries per domain). Agency: Small \$299/mo (25-query pool), Medium \$599/mo (75-query pool), Large \$999/mo (200-query pool), Global \$4,999/mo (2,000-query pool). **Every Agency tier includes unlimited domains.** The Agency track is a parallel product, not a higher rung on one ladder.

The revenue model follows from that. Agency revenue is a function of query-pool tier, not client count. Adding clients does not add revenue directly; it adds monitored queries, and query volume drives the tier. Modelling agency revenue as a per-client multiple would materially overstate it, and we do not do so here.

BYOK architecture means marginal API cost per customer is effectively zero; gross margin at the Agency tier is approximately 90% after payment processing.

Sales motion: one agency channel sales hire from month 3, targeting 20–25 signed agencies in months 3–15. CS/onboarding hire from month 3 supports agency relationships and end-brand activation. Tech hire from month 2 ships the roadmap and SOC 2 readiness.

Conservative scenario

20 signed agencies, blended at Agency Medium (\$599/mo) = **\$12.0K MRR / approximately \$144K ARR** by month 15. Client deployments per agency are irrelevant to this number by design; they drive retention and expansion, not the subscription. Approximately \$150K ARR exiting month 18 including modest direct inbound.

Base case scenario

25 signed agencies blended across Medium and Large (\$599–999/mo, weighted toward Large as monitoring deepens) = **\$20K–25K MRR / approximately \$240K–\$300K ARR** by month 15. Plus a modest direct-to-brand contribution from Starter-tier signups at larger consumer brands, and a first pharma/healthcare Enterprise engagement. **Approximately \$300K–\$375K ARR exiting month 18.**

What the Seed milestone actually requires

Our Seed gate (Section 9) targets **\$250K ARR with at least 60% originating in the agency channel** — \$150K from agencies. At approximately \$12K ARR per converted agency, that is **twelve to thirteen signed agencies**.

Both scenarios above clear it. The conservative case reaches roughly \$144K in channel revenue on 20 agencies at the lowest realistic tier; the base case clears the channel component with margin and adds direct-to-brand on top. The binding constraint is therefore not price per agency and not the size of the market — it is **conversion rate and time-to-first-close on a motion that does not exist yet**. That is the specific risk this round retires, and it is a hiring problem with a known cost rather than an open question about demand.

We would rather state the mechanic than present a model that closes the gap by multiplying a subscription by a client count. Agency revenue scales with query-pool tier; deployments are free by design. Any projection that charges per client deployment is wrong, and we have not built one.

What we are not projecting

Aggressive direct-to-brand sales (which would require a second sales hire we have not budgeted for), enterprise-vertical sales beyond the first pharma/healthcare engagements (regulated sales cycles run 6–12 months), or international expansion (English-language US/UK only in this raise). Upside where a single agency multiplies aggressively — thirty-plus client brands driving a Global-tier subscription — is real but not relied upon.

8. Team and advisors

Founder

Eric Richmond — Founder and Chief Executive Officer. 30+ years as an SEO and search practitioner, beginning when “search” meant directory listings and progressing through every major shift in the discipline: PageRank, mobile-first indexing, structured data, and now AI search. Built CiteMetrix solo from concept to v4.46.0 — the entire WordPress/PHP plugin across 570+ releases, the mobile PWA stack, the Flask admin portal, and the conversational Analyst integration — with no outside capital. Network across pharma, healthcare, and consumer-brand agencies established over three decades in the agency ecosystem.

Operating model post-raise: the tech hire takes over backend and infrastructure ownership; sales and CS hires take over the corresponding GTM motions.

Strategic advisor

Brian Cauley — long-standing strategic advisor. Inputs into GTM positioning, channel strategy, and capital-markets approach. Advisor-only relationship; not a co-founder, not salaried. Equity grant TBD at close, structured as a standard advisor agreement.

Planned hires

Senior backend / AI engineer (~\$160K base loaded, month 2). Scales scan queue infrastructure, delivers roadmap, owns SOC 2 technical readiness. Removes the single-founder dependency on the platform.

Agency channel sales (~\$130K base + \$60K OTE, month 3). Owns agency acquisition, converts the existing pipeline, develops partner-enablement materials. A specific AEO channel operator has been identified for this role. Profile: agency sales experience, ideally with prior SaaS-to-agency-channel work.

Customer success / onboarding lead (~\$95K base loaded, month 3). Combined onboarding and CS role for the first 50 accounts. Owns first-scan briefing calls, the customer health dashboard, and quarterly business reviews at the Agency tier.

All three report to the founder.

9. Use of funds and runway

Eighteen-month runway with the following allocation:

Category	Amount	%	Notes
Tech hire (loaded)	\$360K	21%	~\$160K base + benefits + equity overhead, 18 months
Sales hire (loaded)	\$330K	19%	~\$130K base + \$60K OTE + tools, 18 months
CS / onboarding hire (loaded)	\$180K	10%	~\$95K base loaded, 18 months
Founder salary	\$180K	10%	\$120K/yr — standard pre-seed founder draw
Marketing & channel investment	\$200K	11%	Agency partner enablement, content, conferences
Product & infrastructure	\$150K	9%	API costs, hosting, dev tools, internal CRM
SOC 2 + initial legal	\$80K	5%	SOC 2 Type II audit, Delaware conversion, IP, contracts
Financial services + ongoing legal	\$50K	3%	Bookkeeping, tax prep, outside counsel retainer
Reserve / contingency	\$220K	12%	Buffer for hiring delays, channel opportunities
Total	\$1.75M	100%	

Milestone gates for the next raise

Seed target Q2 2028. Gates: (1) **\$250K+ ARR with at least 60% originating in the agency channel**, (2) SOC 2 Type II certified, (3) first pharma/healthcare Enterprise customer signed, (4) at least one named-brand case study from the agency channel.

Section 7 sets out honestly where the base case lands against gate (1) and what would need to happen to close it.

10. Risks and mitigations

Competitive repositioning

Risk: A funded competitor — Profound is the obvious candidate — repositions to address the brand buyer before we have established the agency channel. They have the capital to hire brand-side sales talent and the recognition to win competitive evaluations.

Mitigation: Speed and channel depth. The agency channel takes months to build credibility and tens of agencies to be defensible; an incumbent cannot replicate it in a quarter. The methodology and transparency position compounds over the same timeframe.

Agency channel adoption

Risk: One agency expanding to nine client brands proves that one agency wants this. It does not prove twenty will. The channel-led thesis fails if agency adoption is slower or shallower than projected — and no agency has yet converted to paid.

Mitigation: The conservative scenario in Section 7 produces a viable trajectory at 20 agencies on the lowest realistic tier. The sales hire is the gating constraint; we have budgeted for a quality hire and are willing to delay if the right candidate is not available in month 3. The anchor relationship is unconverted, and we present it as pipeline rather than revenue.

AI platform API risk

Risk: Citation tracking depends on AI platform APIs remaining accessible and reasonably priced. A platform that closes its API or imposes prohibitive costs would disrupt monitoring of that platform.

Mitigation: BYOK architecture means platform-cost risk is borne by the customer, not by CiteMetrix. Closure risk is mitigated by tracking nine platforms — the diversity is a feature, not a marketing claim — and by the Platform Registry, which lets us add platforms as the landscape shifts. Where direct API access is unavailable (Google AI Overview), we use specialist providers.

Solo founder execution

Risk: Single-founder companies fail more frequently than multi-founder companies. Eric is the founder; Brian is an advisor, not a co-founder. Every line of the platform to date is the founder's.

Mitigation: The three planned hires are explicitly structured to address this: tech owns backend and infrastructure, sales owns the channel, CS owns customer relationships. By month 6 the operating team is four people. Seed milestones include adding a second senior leader as a formal executive.

Category timing

Risk: AI search adoption could plateau before brand-side spending materializes. The category exists and the buyer is materializing; whether brand-side budget materializes at the pace we project is the meaningful timing question.

Mitigation: The agency-channel motion converts *existing* agency budget into CiteMetrix revenue. We are not waiting for brand-side budgets to expand. If they expand on schedule we have upside; if they expand slowly the agency motion still funds operations.

Security and compliance posture

Risk: Enterprise and regulated-industry buyers will ask for SOC 2 Type II. We do not have it. Profound does.

Mitigation: Controls are in place and documented; the third-party audit is funded by this round and is a stated Seed gate. We disclose the status plainly on our public methodology page rather than implying certification — a posture that has been an asset in enterprise conversations rather than a liability.

11. Deal terms summary

Corporate form. The terms below presuppose a Delaware C-corporation with an authorized preferred class. The company currently operates as CiteMetrix LLC (Connecticut); conversion will be completed at or before closing. Founder vesting will be established as part of that conversion.

Raise amount	\$1,750,000
Pre-money valuation	\$5,000,000
Post-money valuation	\$6,750,000
Implied dilution	~25.9%
Security type	Priced equity (Seed Preferred)
Minimum investment	\$25,000 (negotiable)
Use of funds	See Section 9
Runway	18 months from initial closing
Next raise target	Seed, Q2 2028, \$4–6M at \$15–20M pre
Initial closing	31 December 2026 — minimum \$500K
Final closing	31 March 2027 — up to \$1.75M
Pro-rata rights	Standard, available to qualified investors
Information rights	Quarterly investor update, annual financials
Board observer seat	Available to lead investor at \$500K+

A full term sheet is available on request. Structure assumes industry-standard pre-seed terms: 1× non-participating liquidation preference, broad-based weighted-average anti-dilution, single-trigger acceleration on founder shares, and standard founder vesting established at conversion. Final terms subject to lead investor negotiation.

12. Important notices and contact

Confidentiality. This prospectus is confidential and provided solely for the use of the recipient. Recipients may not reproduce, distribute, or share this document with third parties without prior written consent.

No offer to sell securities. This document is for informational purposes only and does not constitute an offer to sell or a solicitation of an offer to buy any securities. Any offer of securities will be made only by a definitive offering document delivered to qualified, accredited investors, and only in jurisdictions where such offer is lawful. CiteMetrix is not registered as an investment adviser or broker-dealer; nothing in this document should be construed as investment advice.

Forward-looking statements. This document contains forward-looking statements about market opportunity, product roadmap, financial projections, and competitive positioning. Forward-looking statements are based on current expectations and assumptions that may prove incorrect. Actual results may differ materially. Pre-seed financial projections in particular should be discounted heavily; the assumptions on which they rest are stated explicitly in Section 7 so investors can apply their own judgment.

Accredited investors only. Any investment will be made under exemption from registration under U.S. securities laws and will be available only to accredited investors as defined in Regulation D under the Securities Act of 1933, as amended.

Contact

Eric Richmond — Founder & CEO · eric@citemetrix.com · (860) 682-9104

Investor relations: investors@citemetrix.com · citemetrix.com/investors

CiteMetrix LLC · Norwalk, CT · citemetrix.com

Prepared with AI assistance; all data, figures, and financial terms independently verified by CiteMetrix LLC.